

IRE Assignment-1, Component-1 Design Note

Lexical and Semantic Retrieval on EB-NeRD and MIND

Name: Naman Singhal **Roll Number:** 2024114013

1 What was built

A config-driven pipeline from raw archives to metric reports in one command (`make data` for Q1's feature store, `make all` for everything) across **EB-NeRD demo** (50,080 impressions, 1,935 distinct users), **EB-NeRD small** (477,534 / 18,827) and **MIND-small** (230,117 / 94,057), counted across all three splits. Both normalise into one schema (`impressions*.parquet`: `id`, `user`, `timestamp`, `candidates[]`, `clicked[]`, `history[]`; `articles.parquet`: `id`, `title`, `abstract`, `body`, `category`, `entities[]`), one code path rather than a Danish branch and an English one. `entities` normalises two source shapes into one `List(String)` (74%/87% non-empty).

Three scorers run over that schema, **bm25**, **emb**, **fused**, on two tracks: **re-rank** (score the pool the log actually showed; leaderboard-comparable) and **retrieval** (ignore the pool, pull the top-200 from the whole corpus, the "candidate generation to a few hundred" the spec asks for). The query for both is the user's most recent clicked titles, `history_len` of them: 30 on EB-NeRD, 100 on MIND, measured per dataset in Section 6. The harness's own metrics are in Section 7.

2 Temporal split

Both datasets ship their own train/test boundary, so I adopted it as `test_start` rather than inventing one, and carved `val` out of the earlier block by time (EB-NeRD small train/val/test: 168,522/64,365/244,647; MIND-small: 95,071/61,894/73,152 impressions). EB-NeRD's day boundary is **07:00**, not midnight, a detail that only shows up in the data.

EB-NeRD ships a separate `history.parquet` per block whose window ends exactly where its impressions begin, so the behaviour-window boundary is given rather than guessed (the alternative, all 7 provided train days with validation halved, buys ~38% more train data at the cost of 7-day-stale test history; nothing here is trained, so fresh history wins).

Enforcement, not discipline. `split.py` asserts on every build, failing rather than warning: strict `train < val < test`, history-precedes-impression, `clicked` a subset of `candidates`, every candidate resolves. Where impossible it says so: MIND's history has no

per-item timestamps, so the guarantee there is structural rather than verified, not passed vacuously.

`tests/test_leakage.py` (`make test`, 19 tests, ~20s) re-checks the same properties from *outside*, reading only `data/processed/`: the build asserts what it meant to write, the tests assert what is on disk, plus two checks that exist only there (`datasets.yaml` cutoffs still match the artifacts; no module outside `eval/run.py` mentions a serving-unavailable column). Verified non-vacuous by injecting a future click and watching it fire.

3 Lexical axis

bm25s over title+abstract with a per-dataset Snowball stemmer (`rank_bm25` scores one query at a time in pure Python, unusable at 65k documents; Pyserini/Lucene scales further but drags in a JVM).

Stemming is not free, and not what I assumed. The brief specifies no preprocessing, so this was my choice. On val, `stem - none` is **+0.0019** [+0.0007, +0.0029] **AUC on MIND, zero on EB-NeRD, and significantly negative on recall@200 for both** (-0.0020, -0.0023): it trades precision for variant-matching, worth it ordering ~8 candidates, harmful picking 200 from 20,738. Kept as right for the headline metric, but a track-dependent trade. Danish compound splitting from corpus vocabulary fired on 15% of tokens and made both **worse**: decompounding needs a lexicon, not statistics.

The finding that mattered. EB-NeRD first scored AUC **0.5035**, random. Rather than report that, I checked whether it was my bug: an **oracle test** (query = the clicked article's own title) ranked it #1 in 300/300 impressions at AUC 0.9990, and a history-length sweep was flat, ruling out misalignment and query dilution. MIND worked at 0.58; that asymmetry located the cause: **bm25s ships no Danish stopword list**, so thirty concatenated Danish titles were mostly function words. The Snowball list moved EB-NeRD to **0.5232**. A missing stopword list is silent, it degrades without erroring.

4 Semantic axis

MIND encodes title+abstract with **all-MiniLM-L6-v2** (**11m16s**, cached once, 65,238-article union across

its two files). User vectors are the mean of the last `history_len` history vectors (100 on MIND, 30 on EB-NeRD; Section 6), L2-normalised; scoring is cosine with FAISS IndexFlatIP (exact at this scale; IVF/PQ deferred to 10× analysis).

The encoder is a measurement, not a preference. Q3 allows the provided article vectors *or* your own from BERT/XLM-RoBERTa, so I ran the identical pipeline on each candidate: same splits, same BM25, same fusion, only the article vectors differ. Selection is on **val**, since selecting on test tunes a choice on the number being reported.

On shipping contrastive_vector: it is an official EB-NeRD artifact from the same `artifacts/` folder as the other two, so it sits in Q3’s “provided article embeddings” arm, word2vec and mBERT are that arm’s named examples, not its full extent. Strongest EB-NeRD encoder measured, by +0.0197 val AUC over my own XLM-R, gap well outside both intervals.

One property predicts all four: **was the model trained so cosine distance means topical similarity?** Contrastive training optimises exactly that and wins; XLM-R paraphrase fine-tuning approximates it and comes second; word2vec (predict neighbouring words) does not; raw mBERT (fill in blanks) emphatically does not, its vectors sitting in a narrow cone where everything looks alike. Size predicts nothing: the 300-d model beats one 768-d model. Fusion detects the mBERT failure unprompted, tuning alpha to 0.00 on demo, discarding the semantic axis outright.

MIND kept MiniLM, and that is the sharper lesson. Against the same XLM-R, MiniLM wins on **val** (0.6338 vs 0.6256) and *loses* on **test** (0.6339 vs 0.6364), a real reversal, near-disjoint intervals. The rule was fixed before looking, so MiniLM stays; “take the better test number” would have flipped it, costing 0.0025 test AUC to obey, the entire point of having a rule. An English specialist beats a 50-language generalist on English, while that generalist beats a 2013 word-vector model on Danish.

A four-way MIND ablation settles the accuracy-vs-cost question. MiniLM (384-d, 0.6356) ties `mpnet` (768-d, 0.6344, not significant) and beats `e5` (0.6067) and `bge` (0.5899) outright, at half the FAISS index size and half the per-query cosine cost: no trade-off to weigh, it wins on both axes. That `bge`, built for retrieval, loses by 0.0457 to a small general-purpose model is the clearest evidence that neither dimensionality nor stated training purpose predicts performance here, only measurement does.

5 Candidate generation: recall@K, and a split-dependent winner

Recall@K over the **full catalogue**, not the pool the log showed, is what measures candidate generation: the share of an impression’s clicked articles found in the top K. Cold-start impressions retrieve nothing and score 0 rather than being excluded, since a retriever that cannot serve a user has failed for that user.

On EB-NeRD the retrieval winner flips between val and test, both directions significant: BM25 wins val, embeddings win test, same code, same K, a sign change with disjoint intervals, not noise. MIND shows no such flip. Section 6 finds the identical pattern in the fusion weight, pointing at the dataset rather than either method: val is a 2-day window against test’s 7, each with its own history window, so the relationship genuinely drifts.

Re-ranking and retrieving stay different jobs, by margin now, not direction. The same contrastive vectors beat BM25 by +0.0289 AUC ordering the pool and by only +0.0030 recall@200 searching all 20,738 articles, a 10× gap for the same encoder and split, evidence for different signals per stage. Absolute recall is low everywhere (2–5%), Section 8 explains why, a property of the corpus, not the retriever.

6 Fusion, an addition the brief did not ask for

Q3 asks only to *compare* lexical and semantic; this third system is mine. BM25 scores and cosines are on different scales, and BM25’s moves with query length, so both are z-normalised **within each candidate pool** and mixed $\alpha \cdot \text{emb} + (1 - \alpha) \cdot \text{bm25}$, α tuned on **val**, applied unchanged to test.

On EB-NeRD fusion wins on val and significantly loses on test, on both bundles: the alpha chosen on val does not generalise across the boundary. This is the same reversal Section 5 finds in retrieval, reached through a different measurement, a property of EB-NeRD, not an artefact of either method.

Consequently **the reported EB-NeRD system is embeddings alone**, and only MIND ships fused. Fusion still earns its place: it is the mechanism that *detected* the instability, and on MIND it is a real if small win. Judged on val alone, the tempting shortcut, I would have shipped fusion on all three and been wrong on two.

history_len was shared across datasets and shouldn’t have been. A sweep on MIND val found the constant (30) truncating a real tail: MIND’s median history is 20, not EB-NeRD’s 258, so 30 discarded signal rather than noise. AUC rose significantly to 100 clicks (bm25 +0.0039, emb +0.0018) then plateaued exactly there. EB-NeRD’s own sweep (Section 3) was

Encoder	Arm	Dim	Val AUC	vs BM25 (0.5205)
Ekstra_Bladet_contrastive_vector (<i>ships</i>)	provided	768	0.5506 [.5480, .5531]	beats
paraphrase-multilingual-mpnet (XLM-R)	own	768	0.5309 [.5282, .5336]	beats
Ekstra_Bladet_word2vec	provided	300	0.5098 [.5073, .5124]	loses
google_bert_base_multilingual_cased	provided	768	0.4857 [.4832, .4880]	below random

Table 1: Semantic-only AUC on ebnerd_small val, everything else held fixed.

recall@200	bm25	emb	emb – bm25
EB-NeRD small, val	0.0214	0.0198	−0.0016 [−.0030, −.0002]
EB-NeRD small, test	0.0247	0.0277	+0.0030 [+0.0022, +0.0039]
MIND-small, test	0.0220	0.0239	+0.0019 [+0.0008, +0.0031]

Table 2: Full-corpus retrieval, small bundles; demo tracked these within noise (Section 8).

flat over the same range, so it stays at 30; `history_len` is now per-dataset in config. This table reflects the corrected MIND value.

A second MIND-only signal, entity overlap, ships alongside emb. MIND’s entities are clean and Wikidata-linked; a candidate’s Jaccard overlap with the user’s recent-history entities, blended onto `emb` ($\alpha \cdot \text{entity} + (1 - \alpha) \cdot \text{emb}$), was selected on val (0.20) and checked once on test: **+0.0022** [+0.0018, +0.0026] **AUC**, holding almost exactly (val +0.0023), unlike the fusion alpha above. Not tried on EB-NeRD, its entities are a different extraction (`ner_clusters`) with different coverage.

7 Evaluation harness

Metrics are computed **per impression**: slicing is a mask over that vector, bootstrapping resamples it over impressions, the approximately-independent unit. Comparisons use a **paired** bootstrap on shared resamples, since two systems on the same impressions have correlated errors, an unpaired comparison of independent intervals is far too conservative.

Cold-start needed a defensible definition. “Empty history” fails on EB-NeRD (zero such users, median 258), so the harness uses a per-dataset bottom-decile threshold plus the true zero-history slice where one exists; MIND’s scores 0.5125 identically across systems, correctly, no history means no ranking invented. Beyond-accuracy over the top-10: category intra-list diversity, novelty as self-information against *train* click shares, coverage, sobering at **5–21%** of the catalogue ever surfaced.

8 Observations

Content signals are weak here, and that is the result, not a failure. MIND reaches 0.6391 (Section 6), within reach of published neural baselines for MIND-small (~ 0.65 – 0.67), unsupervised. EB-NeRD

tops out at 0.5397: pools are much smaller (median 8–9 vs MIND’s 23–26) and drawn from a tight recency window, so candidates are already plausible, little for content to separate.

Simple behavioural features are not the shortcut they look like. Train-click popularity scores *below* random on EB-NeRD test (AUC 0.4498 [0.4485, 0.4510] on small, 0.4742 on demo): the pool the log showed is already popularity-filtered, so within it, being popular is mildly *anti*-predictive.

Demo was representative. Every system landed within 0.002 of its demo value at 10× the users; the extra data bought precision, CIs tightened from ± 0.0040 to ± 0.0013 , deciding fusion.

The retrieval track is weak regardless of scorer (recall@200 of 2–5%): EB-NeRD’s corpus spans 1998–2023 while impressions are one week in May 2023, and the publication filter removes only *future* articles, not stale ones. A recency window on the retrievable corpus is the obvious C2 lever.

9 Anti-gaming and serving availability

`total_inviews/total_pageviews/total_read_time` are lifetime aggregates over the whole collection period, embedding the future, and cannot be computed at serving time. Carried into the corpus **only** so the harness can quantify them; no shipped scorer reads them. Test AUC without/with popularity: EB-NeRD demo 0.5384 $\rightarrow$ 0.5793 (+0.0409 [+0.0380, +0.0437]), EB-NeRD small 0.5380 $\rightarrow$ 0.5797 (+0.0417 [+0.0408, +0.0426]).

One serving-unavailable feature is worth +0.042, still around 1.4× the entire semantic axis (+0.0289 emb over bm25, Section 6), the largest legitimate win here; any comparison quietly including it is not measuring the same system.

Also enforced, not trusted: post-click engagement columns are dropped at the split, and retrieval filters `published_time <= impression_time` (867 EB-

AUC	bm25	emb	fused	α	fused – emb
EB-NeRD small, val	0.5205	0.5506	0.5528	0.70	+ 0.0022 [+0.0010, +0.0033]
EB-NeRD small, test	0.5107	0.5397	0.5380	0.70	– 0.0017 [–0.0023, –0.0011]
MIND-small, test	0.5685	0.6369	0.6381	0.75	+ 0.0012 [+0.0005, +0.0018]

Table 3: Fusion against its components. Demo tracked these within noise (Section 8).

NeRD articles postdate the test window).

10 Where it breaks at $10\times$, measured, not projected

The leaderboard submission made this section measured rather than hypothetical: Codabench scores `ebnerd_testset` at **13.5M impressions, 205,925,868 candidate pairs**, $28\times$ `ebnerd_small`. **Memory is the wall, and it is one specific idiom.** Three separate times the break was `.to_list()` pulling Arrow into boxed Python objects: `np.vstack(col.to_list())` peaks at **5.8 GB to build a 385 MB array**, the same idiom OOM-killed the harness at 10 GB RSS adding `recall@K.explode().to_numpy()` plus Polars set intersections fixed all three, byte-identically.

Streaming works and costs little. `pipeline/submit.py` scores in 200k-impression slices: peak RSS stays **flat against dataset size, not linear**, in **13m30s** (33m28s at the same peak under the earlier XLM-R build), a throughput of $\sim 16.7k$ impressions/s; MIND-large’s 2.4M took 4m38s at 2.0 GB. Slice pushdown reaches the parquet row groups, so an offset-13M slice costs the same as one at 0; without it streaming would be $O(n^2)$. MIND needed a one-off TSV $\rightarrow$ parquet conversion, CSV has no such pushdown and OOM-killed the first attempt.

BM25 is expensive, and I first overstated by how much. `get_scores` is dense over the corpus per query; $\sim 2M$ MIND-large histories $\times 120,961$ articles, which I called impractical from the $\sim 10^{11}$ operation count alone. Timed, it is 1.6ms/query, so **~ 1.6 h**, roughly $7\times$ the embedding run: slow, not a wall. Operation counts are not wall-clock for a sparse vectorised kernel, and I should have measured before asserting. The fix is unchanged: a dense per-query scorer is the wrong shape, WAND-style early termination removes the cost. At `ebnerd_large` (600M) FAISS flat is $O(N\cdot d)$ /query, needing IVF-PQ with `nlist` $\sim \sqrt{N}$.

11 Leaderboard submissions

Both competitions score held-out sets far larger than the splits above (`ebnerd_testset` 13.5M impressions, MIND-large test 2.4M), streamed by

`pipeline/submit.py` (Section 10) with unchanged scoring semantics. **Both entries are embeddings-only, not fused:** on EB-NeRD that *is* the reported system (Section 6); on MIND, BM25 would cost $\sim 1.6h$ for +0.0012 AUC, a deliberate trade reported as a different system.

	Sub.	AUC	MRR	n@5	n@10	Rk
EB-NeRD	895220	.5381	.3521	.3868	.4669	138
MIND	898179	.6503	.3198	.3454	.4010	36

Table 4: Final Codabench submissions (comps. 2469, 13967), organisers’ own output.

Superseded: EB-NeRD 893893 (0.5336, XLM-R encoder) and MIND 892088 (0.6460, `history_len=30`, no entity blend); both improved on by the finals above.

Each score reads against the offline number for the system that produced it. 895220 (post-switch) against offline `ebnerd_demo emb test`, 0.5418: leaderboard -0.0037 , unremarkable, a larger independent population regressing toward the mean. 898179 against offline `emb+entity test`, 0.6391: leaderboard +0.0112, the reassuring direction. (An earlier draft mis-mapped 893893 to the wrong encoder here; dropped rather than repeated.)

MIND’s MRR is the one metric that moved against us, 0.3198 official against 0.3548 offline for the same `emb+entity` system, while AUC and both nDCGs rose, consistently across every MIND submission: a definitional gap, not noise. `eval/metrics.py` takes the *first* clicked article’s reciprocal rank; averaging over *all* clicks scores multi-click impressions lower while leaving AUC and nDCG intact. Unverifiable without the organisers’ scorer, so recorded rather than resolved.

Both validated pre-upload: every line a correct-length rank permutation in test-file row order, plus 40 re-ranked independently in float64. Excluded: 897967 (MIND, 0.5012), a `-limit` smoke test that overwrote the real file. **Evidence:** `leaderboard_screenshots/` and `{ebnerd,mind}_scoring_result/`, both under `deliverable/`. Every measurement quoted here, plus the ablations and failures behind them, is recorded in `docs/NOTES.md`.